

# PREDICTING USED CARS PRICES

## IMPORT NEEDED LIBRARIES

```
In [1]: import pandas as pd # for data processing

import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split # for splitting the dataset

from sklearn.preprocessing import LabelEncoder

from sklearn.linear_model import LinearRegression

from sklearn.ensemble import RandomForestRegressor

from sklearn.metrics import r2_score, mean_squared_error
```

## LOADING THE DATASET

```
In [2]: df = pd.read_excel("C://Users//quays//Desktop//used_car_listings.csv.xlsx")
```

## VIEWING THE FIRST 5 ROWS

```
In [3]: df.head()
```

```
Out[3]:
```

	listing_id	make	model	year	trim	body_type	fuel_type	transmission	mileage
0	2	Nissan	Rogue	2024	LT	Sedan	Hybrid	Automatic	16109
1	3	Hyundai	i20	2018	XLE	Crossover	Petrol	Automatic	173239
2	4	Kia	Sportage	2023	EX	Hatchback	Diesel	CVT	36810
3	5	Kia	Seltos	2020	Trend	Pickup	Diesel	Automatic	87749
4	6	Mercedes-Benz	GLA	2019	Platinum	Crossover	Electric	CVT	60853



# DATA CLEANING AND EXPLORATORY DATA

## 1. CHECKING FOR DUPLICATES

```
In [4]: df.duplicated().sum()
```

```
Out[4]: 0
```

## 2. CHECKING FOR NULL SPACES

```
In [5]: df.isnull().sum()
```

```
Out[5]: listing_id      0
        make           0
        model          0
        year           0
        trim           0
        body_type      0
        fuel_type      0
        transmission   0
        mileage        0
        price          0
        condition      0
        Country        0
        seller_type    0
        dtype: int64
```

## 3. CHECKING COLUMNS AND THEIR DATA TYPES

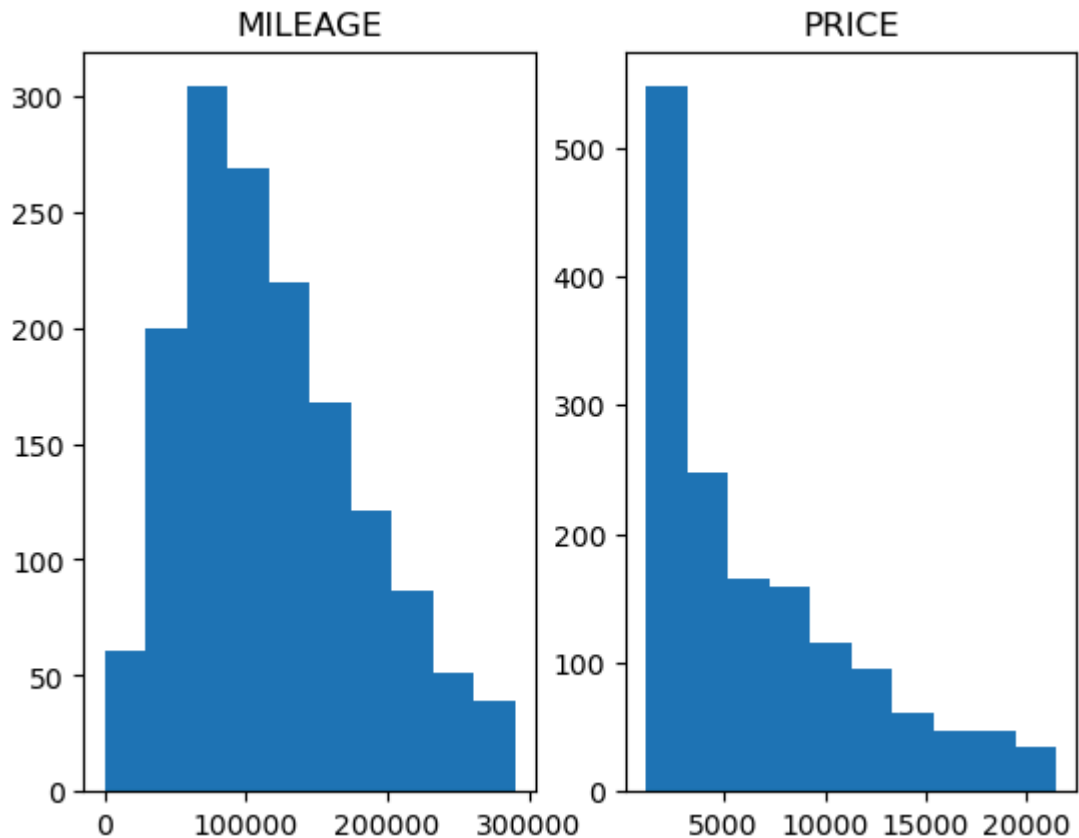
```
In [6]: df.dtypes
```

```
Out[6]: listing_id      int64
        make           object
        model          object
        year           int64
        trim           object
        body_type      object
        fuel_type      object
        transmission   object
        mileage        int64
        price          int64
        condition      object
        Country        object
        seller_type    object
        dtype: object
```

## 4. CAR MILEAGE AND PRICE DISTRIBUTION

```
In [7]: plt.subplot(1, 2, 1)
plt.hist(df['mileage'])
plt.title('MILEAGE')

plt.subplot(1, 2, 2)
plt.hist(df['price'])
plt.title('PRICE')
plt.show()
```



## 5. ENCODING THE STRING COLUMNS TO NUMERICAL VALUES

```
In [8]: encoder = LabelEncoder()
```

```
In [9]: df['make'] = encoder.fit_transform(df['make'])
df['model'] = encoder.fit_transform(df['model'])
df['trim'] = encoder.fit_transform(df['trim'])
df['body_type'] = encoder.fit_transform(df['body_type'])
df['fuel_type'] = encoder.fit_transform(df['fuel_type'])
df['transmission'] = encoder.fit_transform(df['transmission'])
df['condition'] = encoder.fit_transform(df['condition'])
df['Country'] = encoder.fit_transform(df['Country'])
df['seller_type'] = encoder.fit_transform(df['seller_type'])
```

In [10]: `df.head()`

Out[10]:

	listing_id	make	model	year	trim	body_type	fuel_type	transmission	mileage	price	con
0	2	10	44	2024	4	7	3	1	16109	19480	
1	3	5	64	2018	17	2	4	1	173239	4556	
2	4	6	53	2023	1	3	1	2	36810	11536	
3	5	6	46	2020	16	5	1	1	87749	14098	
4	6	9	26	2019	10	2	2	2	60853	17137	

## SPLITTING THE DATA INTO TRAINING AND TESTING

### 1. GROUPING THE DATA INTO DEPENDENT AND INDEPENDENT VARIABLES

In [11]: `df.columns`

Out[11]: Index(['listing\_id', 'make', 'model', 'year', 'trim', 'body\_type', 'fuel\_type',  
'transmission', 'mileage', 'price', 'condition', 'Country',  
'seller\_type'],  
dtype='object')

In [12]: `X = df[['listing_id', 'make', 'model', 'year', 'trim', 'body_type', 'fuel_type',  
'transmission', 'mileage', 'condition', 'Country',  
'seller_type']]  
y = df['price']`

### 2. SPLITTING THE DATA

In [13]: `X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)`

## 1. LINEAR REGRESSION MODEL

In [14]: `lr_model = LinearRegression()`

## TRAINING THE MODEL ON THE DATA

```
In [15]: lr_model.fit(X_train, y_train)
```

```
Out[15]: LinearRegression()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

## MAKING PREDICTIONS

```
In [16]: lr_pred = lr_model.predict(X_test)
```

## MODEL METRICS/PERFORMANCE

```
In [17]: lr_r2 = r2_score(y_test, lr_pred)
lr_r2
```

```
Out[17]: 0.6528444192293608
```

```
In [19]: lr_mse = mean_squared_error(y_test, lr_pred, squared = False)
lr_mse
```

```
Out[19]: 2960.1500399484685
```

## RANDOMFOREST REGRESSOR

```
In [20]: rf_model = RandomForestRegressor()
```

## TRAINING THE MODEL ON THE DATA

```
In [21]: rf_model.fit(X_train, y_train)
```

```
Out[21]: RandomForestRegressor()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

## MAKING PREDICTIONS

```
In [22]: rf_pred = rf_model.predict(X_test)
```

## MODEL METRICS/PERFORMANCE

```
In [23]: rf_r2 = r2_score(y_test, rf_pred)
rf_r2
```

```
Out[23]: 0.8303611888910583
```

```
In [24]: rf_mse = mean_squared_error(y_test, rf_pred, squared = False)
rf_mse
```

```
Out[24]: 2069.2561706265337
```

## CONCLUSION

***THE BEST PERFORMING MODEL IS THE RANDOM FOREST REGRESSOR WHICH SHOWED THE HIGHEST VALUES FOR THE R2 SCORE AND AND THE LOWEST VALUE FOR THE MEAN SCORE ERROR WHICH ARE 0.83 AND 2069.3 RESPECTIVELY.***

```
In [ ]:
```